

College Application Prediction

2023-07-18

PROJECT OVERVIEW

In this exercise, we will predict the number of applications received using the other variables in the College data set.

First we will split the data set into a training set and a test set.

```
library(ISLR2)
```

```
## Warning: package 'ISLR2' was built under R version 4.3.2
```

```
data(College)
head(College)
```

##	Private	Apps	Accept	Enroll	Top10perc	Top25perc
## Abilene Christian University	Yes	1660	1232	721	23	52
## Adelphi University	Yes	2186	1924	512	16	29
## Adrian College	Yes	1428	1097	336	22	50
## Agnes Scott College	Yes	417	349	137	60	89
## Alaska Pacific University	Yes	193	146	55	16	44
## Albertson College	Yes	587	479	158	38	62
##	F.Undergrad	P.Undergrad	Outstate	Room.Board	Books	
## Abilene Christian University	2885		537	7440	3300	450
## Adelphi University	2683		1227	12280	6450	750
## Adrian College	1036		99	11250	3750	400
## Agnes Scott College	510		63	12960	5450	450
## Alaska Pacific University	249		869	7560	4120	800
## Albertson College	678		41	13500	3335	500
##	Personal	PhD	Terminal	S.F.Ratio	perc.alumni	Expend
## Abilene Christian University	2200	70	78	18.1	12	7041
## Adelphi University	1500	29	30	12.2	16	10527
## Adrian College	1165	53	66	12.9	30	8735
## Agnes Scott College	875	92	97	7.7	37	19016
## Alaska Pacific University	1500	76	72	11.9	2	10922
## Albertson College	675	67	73	9.4	11	9727
##	Grad.Rate					
## Abilene Christian University	60					
## Adelphi University	56					
## Adrian College	54					
## Agnes Scott College	59					
## Alaska Pacific University	15					
## Albertson College	55					

```
set.seed(123)
indis <- sample(1:nrow(College), round(2/3*nrow(College)), replace = FALSE)
college_train <- College[indis, ]
college_test <- College[-indis, ]
```

Now I will fit a linear model using least squares on the training set, and report the test error obtained.

```
lm.fit <- lm(Apps~., data = college_train)
lm_pred <- predict(lm.fit, college_test )
summary(lm.fit)
```

```
##
## Call:
## lm(formula = Apps ~ ., data = college_train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -3098.1  -435.7   -32.6    326.9   6524.3
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -320.63000   483.82540  -0.663  0.507830
## PrivateYes  -631.06608   166.38884  -3.793  0.000167 ***
## Accept       1.22765     0.05907  20.782 < 2e-16 ***
## Enroll       0.07342     0.22242   0.330  0.741483
## Top10perc    45.28449     6.30692   7.180 2.54e-12 ***
## Top25perc   -12.88783     5.12008  -2.517 0.012144 *
## F.Undergrad  0.02496     0.04024   0.620  0.535386
## P.Undergrad  0.03394     0.03505   0.968  0.333304
## Outstate    -0.06350     0.02155  -2.947 0.003361 **
## Room.Board   0.20100     0.05392   3.728 0.000215 ***
## Books        0.16346     0.27890   0.586  0.558084
## Personal    -0.03987     0.07418  -0.537  0.591204
## PhD         -6.76818     5.36695  -1.261  0.207866
## Terminal    -5.29390     5.82889  -0.908  0.364201
## S.F.Ratio   -0.13458    14.77294  -0.009  0.992735
## perc.alumni -7.16431     4.68079  -1.531 0.126506
## Expend       0.08032     0.01338   6.005 3.69e-09 ***
## Grad.Rate    9.82319     3.37117   2.914 0.003730 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 980.1 on 500 degrees of freedom
## Multiple R-squared:  0.918, Adjusted R-squared:  0.9153
## F-statistic: 329.5 on 17 and 500 DF, p-value: < 2.2e-16
```

```
Test_error_linear <- mean((college_test$Apps - lm_pred)^2)
Test_error_linear
```

```
## [1] 1684049
```

#The test error is 1684049

Now, we will fit a ridge regression model on the training set, with λ chosen by cross-validation. Report the test error obtained.

```
library(glmnet)
```

```
## Loading required package: Matrix
```

```
## Warning: package 'Matrix' was built under R version 4.3.3
```

```
## Loaded glmnet 4.1-8
```

```
set.seed(123)
X_train = model.matrix(Apps~., data = college_train)
X_test = model.matrix(Apps~., data = college_test)
#Choosing lambda using cross-validation
cv.out = cv.glmnet(X_train, college_train$Apps, alpha=0)
sel = cv.out$lambda.min
sel
```

```
## [1] 311.779
```

```
#fitting ridge model
ridge_mod = glmnet(X_train, college_train$Apps, alpha = 0, lambda=sel)
#Make predictions
ridge_pred = predict(ridge_mod, s=sel, newx = X_test, type = "response")
#Calculate test error
summary(ridge_pred)
```

```
##          s1
## Min.      : -361.0
## 1st Qu.:   861.8
## Median :  1760.8
## Mean     :  3167.3
## 3rd Qu.:  3642.1
## Max.     : 27483.2
```

```
Test_error_ridge <- mean((ridge_pred - college_test$Apps)^2)
Test_error_ridge
```

```
## [1] 2791017
```

#The best lambda by cross validation is 311.779 and the test error is 2791017

Now, we will fit a lasso model on the training set, with λ chosen by crossvalidation. Report the test error obtained, along with the number of non-zero coefficient estimates

```
#first choosing best lambda
set.seed(123)
cv.out_2 = cv.glmnet(X_train, college_train$Apps, alpha=1)
sel2 = cv.out_2$lambda.min
sel2
```

```
## [1] 6.120348
```

```
#Fitting lasso model
lasso_mod = glmnet(X_train, college_train$Apps, alpha=1, lambda=sel2)
#Make predictions
lasso_pred = predict(lasso_mod, s=sel2, newx=X_test)
Test_error_lasso <- mean((lasso_pred -college_test$Apps)^2)
Test_error_lasso
```

```
## [1] 1692748
```

```
coefficient <- predict(lasso_mod, s = sel2, type = "coefficients")

coefficient[coefficient!=0]
```

```
## [1] -409.19194117 -613.21678718 1.22135193 0.09502764 41.68875513
## [6] -9.93420805 0.02373219 0.02791910 -0.05716962 0.18949382
## [11] 0.11856415 -0.02360973 -6.01752912 -5.11448728 -6.85847164
## [16] 0.07907405 8.88004892
```

```
which(coefficient!=0)
```

```
## [1] 1 3 4 5 6 7 8 9 10 11 12 13 14 15 17 18 19
```

```
numberofnonzero <- sum(coef(lasso_mod, s = sel2) != 0)
numberofnonzero
```

```
## [1] 17
```

#The best lambda by cross validation is 6.120348, the test error is 1692748 and the non-zero coefficient estimates are also listed accordingly

Now we will fit a PCR model on the training set, with M chosen by crossvalidation. Report the test error obtained, along with the value of M selected by cross-validation.

```
library(pls)
```

```
## Warning: package 'pls' was built under R version 4.3.3
```

```
##
## Attaching package: 'pls'
```


```
## The following object is masked from 'package:stats':
##
## loadings
```

```
set.seed(123)

pcrfit <- pcr(Apps~., data=college_train, scale=TRUE, validation="CV")

validationplot(pcrfit, val.type = "MSEP")
```

Apps



```
summary(pcrfit)
```

```
## Data:      X dimension: 518 17
## Y dimension: 518 1
## Fit method: svdpc
## Number of components considered: 17
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps  6 comps
## CV              3370    3336    1680    1631    1363    1303    1257
## adjCV           3370    3336    1678    1630    1357    1299    1253
##      7 comps  8 comps  9 comps 10 comps 11 comps 12 comps 13 comps
## CV          1202    1201    1169    1169    1168    1174    1176
## adjCV        1195    1196    1166    1167    1165    1171    1173
##      14 comps 15 comps 16 comps 17 comps
## CV           1176    1176    1029    1029
## adjCV         1173    1173    1025    1025
##
## TRAINING: % variance explained
##      1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps  8 comps
## X          31.765   57.84   64.68   70.19   75.49   80.39   84.01   87.40
## Apps        3.386   75.80   77.45   84.75   86.02   86.91   88.03   88.22
##      9 comps 10 comps 11 comps 12 comps 13 comps 14 comps 15 comps
## X          90.57   93.02   95.07   96.93   98.02   98.88   99.40
## Apps        88.84   88.89   88.94   88.98   89.03   89.03   89.23
##      16 comps 17 comps
## X          99.82   100.0
## Apps        91.74   91.8
```

#The lowest MSEP occurs at around M = 17 which can be confirmed from the summary as well ignoring the rest of the components which we neglect for now as they do not solve the purpose.

```
#predicting using M = 17 found by cross validation
pcrfit1 <- pcr(Apps~., data=college_train, scale=TRUE, ncomp=17)
prediction <- predict(pcrfit1, college_test, ncomp=17)


#test error
Test_error_pcr <- mean((prediction-college_test$Apps)^2)
Test_error_pcr
```

```
## [1] 1684049
```

#also confirmed the M value by changing the value of ncomp value from 8 to 17 and got the minimum value of test error at 17. Hence, considered M = 17 in the final answer. #The test error is 1684049

Now we will fit a PLS model on the training set, with M chosen by crossvalidation. Report the test error obtained, along with the value of M selected by cross-validation.

```
set.seed(123)
#Fit and determine M based on CV results
plsfit = pls(Apps~., data=college_train, scale=TRUE, validation="CV")
validationplot(plsfit, val.type = "MSEP")
```



```
summary(plsfit)
```

```
## Data:      X dimension: 518 17
## Y dimension: 518 1
## Fit method: kernelpls
## Number of components considered: 17
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps  6 comps
## CV           3370    1513    1233    1138    1121    1099    1045
## adjCV        3370    1511    1236    1136    1117    1092    1040
##      7 comps  8 comps  9 comps 10 comps 11 comps 12 comps 13 comps
## CV           1031    1028    1029    1030    1027    1028    1028
## adjCV        1027    1025    1026    1026    1024    1025    1025
##      14 comps 15 comps 16 comps 17 comps
## CV           1029    1029    1029    1029
## adjCV        1025    1025    1025    1025
##
## TRAINING: % variance explained
##      1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps  8 comps
## X          26.30   42.01   63.26   67.75   71.41   74.08   77.53   80.83
## Apps       80.53   86.92   89.34   90.16   91.05   91.71   91.77   91.79
##      9 comps 10 comps 11 comps 12 comps 13 comps 14 comps 15 comps
## X          83.35   86.14   89.53   91.21   93.22   94.67   97.06
## Apps       91.79   91.80   91.80   91.80   91.80   91.80   91.80
##      16 comps 17 comps
## X          99.11   100.0
## Apps       91.80   91.8
```

#From the summary and the plot, the lowest MSEP occur at M = 17.

```
#making prediction with M = 17
plsfit1 = plsr(Apps~., data=college_train, scale=TRUE, ncomp=17)
prediction = predict(plsfit1, college_test, ncomp = 17)
#test error
Test_error_pls <- mean((prediction - college_test$Apps)^2)
Test_error_pls
```

```
## [1] 1684049
```

#also confirmed the M value by changing the value of ncomp value from 8 to 17 and got the minimum value of test error at 17. Hence, considered M = 17 in the final answer. #the test error is 1684049

Finally explaining the results obtained. We will check how accurately can we predict the number of college applications received. Also, is there much difference among the test errors resulting from these five approaches?

```
Test_error_linear
```

```
## [1] 1684049
```

```
Test_error_ridge
```

```
## [1] 2791017
```

```
Test_error_lasso
```

```
## [1] 1692748
```

```
Test_error_pcr
```

```
## [1] 1684049
```

```
Test_error_pls
```

```
## [1] 1684049
```

We see that the test errors for linear regression, PCR and PLS are relatively close, the test errors of ridge regression is a little higher. Lasso also has a little more value as compared to linear, pcr and pls. There are not much differences in the test errors except for that in ridge regression. We can predict the number of college application received with reasonable accuracy.

Data Preparation

I began by splitting the dataset into a training set and a test set. This was done to ensure that the models could be evaluated on unseen data, helping to prevent overfitting and giving a better estimate of their performance in real-world scenarios.

Linear Regression

The first model I used was linear regression, which served as a baseline for understanding the relationship between the number of applications and the other variables in the dataset. By fitting this model to the training data, I could predict the number of applications in the test set and calculate the associated test error. This provided a straightforward way to gauge how well the other variables explained the variation in application numbers.

Ridge Regression

Next, I explored Ridge Regression, a technique designed to address issues of multicollinearity by regularizing the coefficients of the model. This helps to prevent overfitting, especially when dealing with datasets that have highly correlated variables. I used cross-validation to select the optimal value for the regularization parameter, λ . Ridge Regression slightly increased the test error compared to linear regression, indicating that while it may help with multicollinearity, it didn't drastically improve predictive accuracy for this particular dataset.

Lasso Regression

I then applied Lasso Regression, which not only regularizes the coefficients like Ridge Regression but also performs variable selection by shrinking some coefficients to zero. This makes Lasso particularly useful when dealing with high-dimensional data, as it can simplify the model by eliminating less important variables. After selecting the optimal λ via cross-validation, I found that the test error was similar to that of linear regression, but with fewer variables contributing to the prediction. This indicated which predictors were most influential, offering insights into the key factors driving college applications.

Principal Component Regression (PCR)

Principal Component Regression (PCR) was the next technique I used. PCR reduces the dimensionality of the data by transforming the original variables into a set of uncorrelated components before fitting the regression model. I used cross-validation to determine the optimal number of components (M) to include. The results showed that using 17 components provided the lowest test error, which was comparable to the linear regression model. This indicated that while PCR effectively reduced the dimensionality, it did not significantly improve the predictive accuracy.

Partial Least Squares (PLS)

Finally, I applied Partial Least Squares (PLS), which, like PCR, is a dimensionality reduction technique. However, PLS differs by considering the relationship between the predictors and the response variable when forming the components. Cross-validation also suggested using 17 components, resulting in a test error similar to that of PCR and linear regression. This reinforced the idea that while dimensionality reduction can be useful, it didn't offer a substantial performance boost for this dataset.

Results and Conclusions

After comparing the test errors from all the models, I found that the differences were relatively small. Linear regression, PCR, and PLS produced similar test errors, while Ridge Regression had a slightly higher error. Lasso Regression performed similarly to linear regression but with fewer variables, providing a more interpretable model.

These results suggest that for this dataset, simpler models like linear regression are sufficient to achieve reasonable predictive accuracy. More complex models like Ridge and Lasso Regression or dimensionality reduction techniques like PCR and PLS did not significantly outperform linear regression. This indicates that the relationships in the data were relatively straightforward, and more advanced methods were not necessary for accurate predictions in this case.